

Compiler-driven SIMD offloading using Arm Streaming SVE

Mohamed Husain Noor Mohamed, Adarsh Patil,
Latchesar Ionkov, Eric Van Hensbergen

5th Workshop on Compiler-assisted Correctness Checking and Performance Optimization for HPC (C3PO)
13th June 2025

Outline

- Introduction
 - Arm SIMD instructions and Streaming SVE
- Motivation
 - Compiler auto-vectorizers limitations
- Architectural considerations for code-gen & cost-model
 - Synchronizations
 - Memory access characteristics
- Compiler extension proposals
 - LLVM
 - MLIR
- Conclusions and takeaways

Outline

- Introduction
 - Arm SIMD instructions and Streaming SVE
- Motivation
 - Compiler auto-vectorizers limitations
- Architectural considerations for code-gen & cost-model
 - Synchronizations
 - Memory access characteristics
- Compiler extension proposals
 - LLVM
 - MLIR
- Conclusions and takeaways

Arm vector ISA extensions

SME

Vector length agnostic (VLA)
Implementation dependent (128-2048 bits)

Streaming SVE

Matrix outer product of
vectors

Load, store, insert &
extract matrix vectors

Matrix transposition

SVE2

NEON DSP++

Multi-precision arithmetic

Match detect & histogram

Non-temporal G/S

Bitwise ternary logic

Bitwise permute

AES, SHA3, SM4 crypto

SVE

Scalable
vectors

Per-lane
predication

Gather-load
Scatter-store

Speculative
vectorization

ML extension:
(FP16 + DOT
product)

Arm v8.6
BF16, FP &
Int8 matmul

NEON

Fixed 128-bit vector length

Basic
conditional
select

Int, FP,
saturating
ops


Zip/unzip,
extract, insert

Bitwise and,
shifts, masks

Structured
loads/stores

Streaming SVE deep-dive

- **Streaming SVE:** High-throughput processing of large datasets with simple loops/minimal branching
 - Can execute a subset of SVE2 instructions with vector length = matrix tile width
 - Some SVE and most NEON instructions become *illegal*
 - All other AArch64 scalar integer, FP, and control flow instructions *remain legal*
- **SMSTART & SMSTOP** instructions to **enter/exit Streaming SVE mode**.
 - Cheap write to **PSTATE.SM** execution mode bit, tracked at core front-end
 - No need to call OS/Hyp to switch, registers are cleared upon mode exit
- **Streaming SVE architecture**
 - **SSVE** unit* is disaggregated from core but architecturally acts as part of each core
 - Instructions sent to SSVE unit when streaming mode enabled; dedicated bus for control path – addr, offsets etc.
 - Data path: uses load/stores; private L1, high-bandwidth connection to shared L3
 - Multiple implementations possible from the spec



Key question: What's the best way to use Streaming SVE?

- Intrinsics (as used in Geekbench)
- Libraries (e.g., Arm Compute Library)
- Compilers ← **Today's talk!**

Goal: Enable performant, precise, repeatable code-gen using auto-vectorizers

Outline

- Introduction
 - Arm SIMD instructions and Streaming SVE
- **Motivation**
 - Compiler auto-vectorizers limitations
- Architectural considerations for code-gen & cost-model
 - Synchronizations
 - Memory access characteristics
- Compiler extension proposals
 - LLVM
 - MLIR
- Conclusions and takeaways

Generating SSVE instructions using auto-vectorizers

1. Streaming mode switch using AArch64 SME Attribute

- Function level annotation `void s_callee(void) __arm_streaming;`
- clang manages PSTATE.SM automatically

2. Force auto-vectorization for streaming-enabled functions using flags

`-mllvm -enable-scalable-autovec-in-streaming-mode`

3. Generate IR with scalable vector type *

○ LLVM

```
define float @add_f32(<vscale x 8 x float> %a, <vscale x 4 x float> %b) {  
    %r1 = call @llvm.vector.reduce.fadd.f32.nxv8f32(float - 0.0, <vscale x 8 x float> %a)  
    %r2 = call @llvm.vector.reduce.fadd.f32.nxv4f32(float - 0.0, <vscale x 4 x float> %b)  
    %r = fadd %r1, %r2  
    ret float %r  
}
```

○ MLIR

```
llvm.func @vector_splat_1d_scalable()->vector<[4] xf32> {  
    %0 = llvm.mlir.constant(dense<0.000000e+00> : vector<[4]xf32>) : vector<[4]xf32>  
    llvm.return %0 : vector<[4]xf32>  
}
```

`llvm.vscale` or `vector.vscale` is unknown, use flags `-mcpu/-mtune` to specify

Issues with auto-vectorizer generated code

- **Correctness**

Generates SVE2 instructions but SSVE != SVE2, **SSVE subset of SVE2** (leads to EXC_BAD_INSTR)
Excluded instructions **depends on implementation** (query architectural flags)

- **Performance**

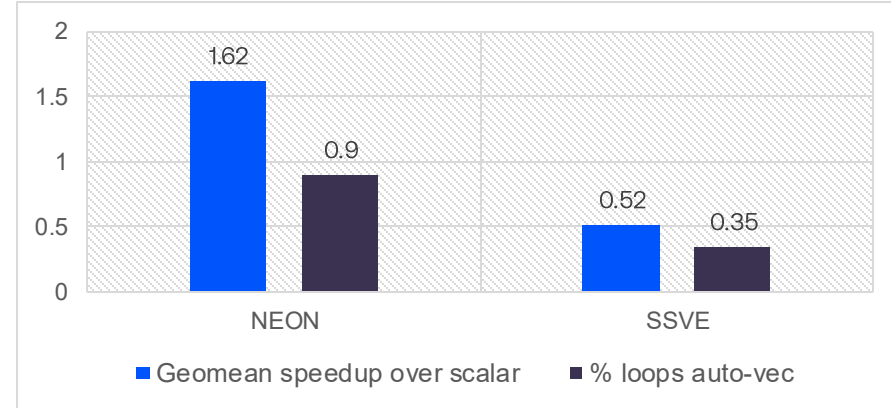
Incompatible cost-model to decide profitability of SSVE vectorization
Why? Different architecture
SVE2 in-core SIMD != SSVE **core-adjacent SIMD offload**

Let's benchmark it!


Platform A, Armv9.2 with SME support
128-bit NEON, 512-bit SSVE unit

Performance of auto-vectorized code

- TSVC_2 (1M iterations, LEN_1D 32000, LEN_2D 256)
Correctness: Only 35% loops are auto-vec to SSVE
Performance: slowdowns over NEON and scalar



- Mandelbrot set (varied input sizes)
SSVE shows slowdown over NEON and scalar



- SPEC 2017
505.mcf_r: 1.9x slowdown over NEON

Outline

- Introduction
 - Arm SIMD instructions and Streaming SVE
- Motivation
 - Compiler auto-vectorizers limitations
- Architectural considerations for code-gen & cost-model
 - Synchronizations
 - Memory access characteristics
- Compiler extension proposals
 - LLVM
 - MLIR
- Conclusions and takeaways

Synchronizations

- GPR and FPR sync
 - **Issue:** Stalls occur when SSVE and core access the same stack cache lines.
 - Cause: SSVE has a private L1 and accesses primarily FP (FPR) stack objects.
 - **Mitigation:** Use padding between GPR- and FPR-accessed stack objects.
Use `--aarch64-stack-hazard-size` flag
 - Effect: Padding moves FPR-only data inward, GPR-only data outward, reducing conflicts.
 - Limitation: Padding can't isolate all objects (e.g., VLAs, shared data).
 - **Impact:** Residual stall penalty increases with loop count — from 17% up to 61%*
- Predicates sync
 - Predicate registers used to select active lanes in SSVE
 - SSVE unit can produce predicates independently
 - **Issue:** Core using SSVE-produced predicates requires synchronization.
Predicate synchronization adds additional latency
 - **Mitigation:** use vector registers instead, place predicate usages as far as possible
 - **Impact:** Stall penalty can be tens of cycles *

*Data gathered by micro-benchmarking on Platform A

Memory access characteristics

- Load-store region table (LSRT) Hazards
 - LSRT manages and synchronizes translation between core and SSVE unit (4KB pages).
 - **Issue:** Hazard condition if core and SSVE access the same data, with one being a write
 - LSRT tracks hazards at aligned 1KB mem regions (4 hazard granules per addr. trans. granule)
 - Effect: New accesses stalled until prior writes in the same region are complete and coherent
 - **Mitigation:**
 - Avoid mixing vector and scalar data in the same data structures, including the stack.
 - Align stack regions to 1KB if data is accessed by both core and SSVE unit
 - Generate spills and fills of both scalar/vector data accessing the same region
 - **Impact:**
 - QPSK modulation kernel in Arm RAL shows 6x slowdown over NEON due to LSRT hazard
 - Using above techniques improves performance by 4x

Memory access characteristics

- Offload overheads
 - Non-linear scaling: SSVE outperforms NEON after input crosses a “performance knee”
 - **Mitigation:** Use conservative vectorization factor (VF) thresholds, default to NEON
- Prefetcher on SSVE unit
 - Optimized for strided and nested access; struggles with arbitrary patterns.
 - **Mitigation:** Fall back to NEON if data-dependent addresses or pointers
- Converting gather/scatter
 - Regular strides can be converted to load/store using ZIP/UNZIP instructions
 - **Mitigation:** Fall back to NEON for random strides

Outline

- Introduction
 - Arm SIMD instructions and Streaming SVE
- Motivation
 - Compiler auto-vectorizers limitations
- Architectural considerations for code-gen & cost-model
 - Synchronizations
 - Memory access characteristics
- **Compiler extension proposals**
 - LLVM
 - MLIR
- Conclusions and takeaways

Cursory thoughts on compiler extensions

- LLVM: Loop Vectorizer pass

- Cost modelling for basic auto-vectorization

Vectorization Factor (VF) for profitability analysis expressed as `<vscale X N X type>`

Code generated compatible for any vscale between `vscale_range(min[,max])`

Issue: Model **assumes linear scaling of performance** with increasing vector lengths – incorrect!

Fixed instruction costs in compiler table

Issue: SVE2 instruction **cost changes in streaming vs non-streaming mode**

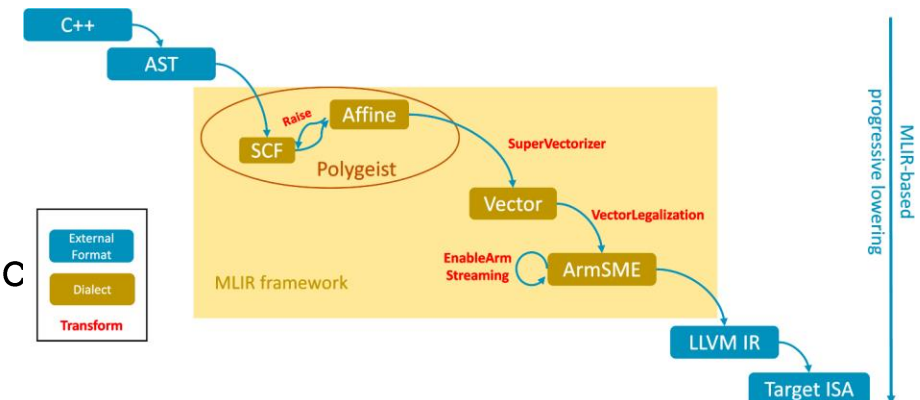
- Vectorization Plans (VPlan) for advanced vectorization strategies

Cost computed by recusing into VPBasicBlocks and into VPRecipies

Issue: Does not model streaming mode overheads (**Instantiate VPlan with static cost for SSVE blocks**)

- MLIR: Polygeist – C/C++ compilation flow

- Vector dialect supports representing Scalable Vectors
- Extend *ArmSME* dialect with SSVE operations – exists
- *EnableArmStreaming* transform to emit smstart/smstop – exists
- **Issue:** *SuperVectorizer* transform does not generate Scalable Vec
requires a new cost model!



Outline

- Introduction
 - Arm SIMD instructions and Streaming SVE
- Motivation
 - Compiler auto-vectorizers limitations
- Architectural considerations for code-gen & cost-model
 - Synchronizations
 - Memory access characteristics
- Compiler extension proposals
 - LLVM
 - MLIR
- Conclusions and takeaways

Takeaways

- New **“offload” SIMD compute** paradigm – Arm Streaming SVE
 - Tightly-coupled, core-adjacent accelerator
 - Ultra-wider vector widths
 - Supports (almost) same ISA as the core
- **Compilers well-suited to enable** SIMD offload
 - Automatic/transparent offload when profitable
 - Reduce programming complexity
 - Generalize and increase applicability
- **Call to action:** Updating compilers to support SIMD offload
 - Auto-vectorizer cost-model – updates to vectorization models, candidates, heuristics
 - Rethink instruction costs – mode and implementation dependent
 - Non-linear performance scaling – offload overheads

arm

Merci

Danke

Gracias

Grazie

谢谢

ありがとう

Asante

Thank You

감사합니다

धन्यवाद

Kiitos

شكراً

ধন্যবাদ

תודה

ధన్యవాదములు

Köszönöm



The Arm trademarks featured in this presentation are registered trademarks or trademarks of Arm Limited (or its subsidiaries) in the US and/or elsewhere. All rights reserved. All other marks featured may be trademarks of their respective owners.

www.arm.com/company/policies/trademarks